

Module 3

U of A: <https://responsibleai.arizona.edu/ai-arizona/ai-guidelines-principles>

Task 3 is the **Academic Integrity Scenario Analysis** — participants apply a disclosure reasoning framework to four deliberately ambiguous AI use situations in a university context, recording a verdict on appropriateness, disclosure, audience, and policy reasoning.

University of Arizona Responsible AI

- [AI Guidelines and Principles](#)
- [AI Tools & Data Use Guide](#)
- [Generative AI in Research](#)

Scenario 1 — **Personalized Student Feedback**

<https://github.com/UA-AI2S/AI-Champions/blob/main/docs/AcademicIntegrityScenarios.md>

Task 3 is the **Academic Integrity Scenario Analysis** — participants apply a disclosure reasoning framework to four deliberately ambiguous AI use situations in a university context, recording a verdict on appropriateness, disclosure, audience, and policy reasoning.

University of Arizona Responsible AI

AI Guidelines and Principles

AI Tools & Data Use Guide

Generative AI in Research

Scenario 1 — **Personalized Student Feedback**

A professor teaching an introductory course with 120 enrolled students uses an AI tool to draft first-pass written feedback on each student's essay. She inputs only a 3-sentence summary of the student's argument and their first name — no grades, no student ID numbers. She spends roughly two minutes reviewing and editing each AI-generated comment before posting it in the course management system under her name. Students receive what feels like personal, thoughtful commentary from their professor. She does not mention AI involvement to anyone.

What makes it contested: The data used is low-sensitivity and no personal records were shared. The professor reviewed and edited every comment. However, students may reasonably assume that feedback signed by their professor was written by their professor —

and the feedback relationship is considered a core part of pedagogy. Does editing AI output constitute "writing" the feedback? Who, if anyone, is owed disclosure — the students, the department chair, or no one?

Policy pressure points: AI transparency; FERPA (marginal here); professional ethics norms around the student-instructor relationship.

Scenario 2 — The Department Newsletter

An office coordinator is responsible for producing the monthly department newsletter. She uses a free AI chatbot to draft it each month, pasting in a list of upcoming events, faculty names, and recent achievements — all publicly available information — and asking the tool to write it in a warm, collegial tone. She edits it lightly, then sends it from the official department email account. She has mentioned this to no one. The newsletter goes to faculty, staff, students, and community partners.

What makes it contested: No sensitive data was used. The content was factually reviewed. The output is routine institutional communication. Yet the newsletter is sent under the department's institutional identity and implies a human voice. Would recipients — or the department chair — evaluate this communication differently if they knew it was AI-drafted? Is disclosure owed, and if so, to whom: the chair, the recipients, or neither?

Policy pressure points: Institutional AI use guidance (if any); professional communication norms; the principle that substantive AI assistance warrants transparency when a reasonable person would find it relevant.

Scenario 3 — The Promotion and Tenure Reference Letter

A tenured senior professor is asked to write a reference letter for a junior colleague's promotion case — one of the most consequential documents in academic career advancement. Drawing on notes he has kept about the colleague's work over several years, he uses an AI tool to draft the letter, then edits it substantially and signs it on university letterhead. The letter is submitted to the tenure committee. Neither the committee, the junior colleague, nor the department chair knows AI was involved. The professor considers the final letter fully representative of his genuine assessment.

What makes it contested: The letter bears his signature and implies personal testimony and professional judgment built over years. He did review and substantially edit it. But the tenure committee is making a high-stakes institutional decision partly on the basis of what they believe is a personally authored document. Is this categorically different from asking a staff member to draft a letter from your notes, which is common practice? Or does the implied personal authorship of a reference letter make AI involvement a material fact?

Policy pressure points: institutional tenure and promotion policy; professional ethics norms around letters of reference; the distinction between "ethically questionable" and "policy-violating."

Scenario 4 — Student Disability Accommodation Letters

An academic advisor in the disability services office is responsible for issuing individualized accommodation letters to students with documented disabilities each semester. Facing a backlog of 45 letters, she uses an AI tool to draft each one, inputting each student's name and their specific accommodation list drawn directly from official disability records. She reviews each letter before sending it from her institutional email. The letters are accurate and the students receive them faster than usual. The AI tool she used has not been reviewed or approved by her institution's IT office.

What makes it contested: This is the highest-stakes scenario on data sensitivity. Student disability status is explicitly protected under FERPA, and the data was entered into a tool without confirmed institutional approval — regardless of how carefully the output was reviewed. Yet the intent was to serve students better and faster, the letters were accurate, and no data was shared externally beyond the AI tool. Did the efficiency benefit justify the data handling? Does the lack of an institutional policy on this tool mean it was implicitly permitted, or does FERPA apply regardless?

Policy pressure points: FERPA (directly applicable); institutional IT data governance policy; the workbook's principle that **"absence of policy is not permission."**

Facilitation note: All four scenarios are deliberately structured so that "Depends" is a defensible verdict for every row. The learning goal is not to reach the "right" answer but to practice citing specific policies, distinguishing ethical concern from policy violation, and identifying who has a legitimate interest in knowing AI was involved.

Scenarios 1 and 2 are the most likely to produce divergent peer verdicts — which makes them the richest for discussion.


Scenario 3 is the most contested on authorship norms;

Scenario 4 is the most clear-cut on data sensitivity;

Module 3: AI for Productivity

Carlos Lizarraga-Celaya edited this page 12 hours ago · [14 revisions](#)

AI for Productivity — Practical Integration in Academic Work



Information	Description
Duration	1 h synchronous + 3 h independent = 4 h total
Learning Purpose	professional role. You will audit your own workflows, complete a real task using AI with full documentation, analyze academic integrity scenarios, and sketch a personal AI workflow diagram — building the foundations of safe, transparent, institutional AI practice.

Before You Begin — Pre-Module Checklist

- Identify your target real work task for Task 2 before the session — a low-to-medium- stakes document or communication task.
- Confirm your [institution's current AI use guidance](#) — have the URL or document title ready.
- Download or bookmark a diagramming tool for Task 4's workflow sketch (same tools as Module 1).
- Have the Module 3 [Scenario Set from the course platform](#) — four scenarios for Task 3.

Learning Objectives for This Module

1. Audit five recurring professional tasks for AI integration potential using a structured multi-dimensional matrix.
2. Complete one real work task using AI with full documentation of prompts, outputs, edits, time, and quality.
3. Apply a disclosure reasoning framework to four academic integrity scenarios using institutional and other academic policy references.
4. Design a personal AI workflow for one task that includes input classification, prompt strategy, verification checkpoint, and disclosure decision.
5. Interpret the Noy & Zhang (2023) and Dell'Acqua et al. (2023) findings in the context of your own workflow documentation.

Reading Guide

Complete readings BEFORE attempting tasks. Annotate as you read — your notes feed directly into the tasks below.


Click here to expand the text box

Reading 1: Experimental Evidence on the Productivity Effects of Generative AI (Noy & Zhang, 2023)

Click to expand a summary NotebookLM Infographic for this reading

Revision Notes: The MIT ChatGPT Productivity Study (Noy & Zhang, 2023)

★ **Goal:** Visual Summary of Generative AI on Professional Writing (Productivity, Quality, Sentiment) 🧠



	Description
Citation:	Noy, S., & Zhang, W. (2023). Experimental evidence on the productivity effects of generative AI. Science, 381(6654), 187–192. NBER Working Paper No. 31161.
Assigned section	Full paper — (has 6 pages in the closed Science version, and the free open version has 15 pages). Read entirely.
Estimated reading time	~20 minutes
Why this matters	This is the highest-quality experimental evidence available on AI productivity effects in professional writing tasks. A randomized controlled trial with 453 college-educated workers found a 37% reduction in task completion time and improved quality, concentrated among mid-skill workers. It is the empirical anchor for Task 2's workflow documentation.

What to look for as you read:

- The experimental design: who were the participants, what tasks did they perform, how was 'quality' measured?
- The 37% time reduction and quality improvement findings — but also WHICH workers benefited most and why.
- The finding that lower-skilled workers benefited MORE from AI access than higher-skilled workers — what does this suggest about the distributional effects of AI?
- The limitations the authors themselves acknowledge — what can and cannot be generalized from this study to your work context?

Guided reading questions (annotate answers as you go):

- **Q1:** What type of tasks were used in the Noy & Zhang experiment? Are these tasks representative of the most cognitively demanding work in your professional role?
- **Q2:** The study found that AI 'compressed' the skill distribution — lower-skilled workers improved more. What are the institutional implications of this finding for how AI tools are deployed in higher education staff contexts?
- **Q3:** The paper found quality improvements in addition to time savings. What metric was used to assess 'quality,' and is that metric applicable to the documents you produce?
- **Q4:** What is one methodological limitation of this study that would affect your confidence in applying its findings to your specific work context?


Key terms — write your own definition after reading:

Term	Your definition (In your own words)
Randomized controlled trial (RCT)	High level of research design where there is a control group (stand or care with no intervention) and an experimental / intervention group. Participants are randomly selected to reduce potential bias and to determine in the intervention actually produced the results of the study.
Treatment effect	A calculation that specifies if there is a true difference in outcomes between the SOC group and the intervention group. The difference can deduce whether there is a causal relationship and the strength of that relationship between groups.
Skill compression	The ability of AI to assist in compressing highly specialized tasks into quickly accessible and usable forms
Task completion time	The amount of time it takes to complete a designated task.
Output quality assessment	Human graders assess creativity, analytical thinking, writing proficiency, and persuasiveness of AI output

Reading 2: [Navigating the Jagged Technological Frontier](#) (Dell'Acqua et al., 2023)

Click to expand a summary NotebookLM Infographic for this reading

Revision Notes: Navigating the Jagged AI Frontier



	Description
Citation:	Dell'Acqua, F., McFowland III, E., Mollick, E. R., Lifshitz-Assaf, H., Kellogg, K., Rajendran, S., ... & Lakhani, K. R. (2023). Navigating the jagged technological frontier: Field experimental evidence of the effects of AI on knowledge worker productivity and quality. Harvard business school technology & operations mgt. Unit working paper, (24-013).
Assigned section	Abstract, Section 1 (Introduction), Section 2 (Experimental Design), Section 4 (Results), and Section 6 (Discussion) (~20 pages total).
Estimated reading time	~25 minutes
Why this matters	The 'jagged frontier' is the most operationally useful mental model for deciding WHEN to use AI assistance. With 750 BCG consultants, this field experiment found +40 performance inside AI's capability frontier - and

	Description
	DEGRADED performance outside it. Understanding what the frontier is jagged, not smooth, prevents over-reliance

What to look for as you read:

- The 'jagged frontier' definition: AI performs well on tasks INSIDE its frontier, poorly on tasks OUTSIDE — but the boundary is irregular and non-obvious.
- The 'centaur' vs. 'cyborg' collaboration styles identified in Section 5 — which best describes how you used AI in your Module 2 tasks?
- The inside-the-frontier vs. outside-the-frontier performance differences — what tasks in YOUR professional context sit inside vs. outside?
- The authors' discussion of over-reliance: why did consultants who used AI on outside- frontier tasks perform WORSE than those who didn't use AI?

Guided reading questions (annotate answers as you go):

- **Q1:** What is the 'jagged frontier' and why is it 'jagged' rather than a smooth capability boundary?
- **Q2:** How did Dell'Acqua et al. identify which tasks were 'inside' vs. 'outside' the frontier experimentally? What implications does this classification method have for practitioners who cannot run experiments?
- **Q3:** What does the 'centaur' collaboration style mean in this paper? What does the 'cyborg' style mean? Which is more appropriate for the task you will complete in Task 2?
- **Q4:** The paper found performance degradation for AI-assisted workers on outside- frontier tasks. What cognitive mechanism do the authors propose to explain this? Do you find the explanation convincing?

Key terms — write your own definition after reading:

Term	Your definition (In your own words)
Jagged technological frontier	AI has an uneven capacity to perform well. Output may be accurate for complex tasks but can fail in simple tasks. It is hard to know the boundaries.
Centaur collaboration	An AI user that splits work into AI and human (self) work
Cyborg collaboration	An AI user that integrates AI onto all of their task flow
Over-reliance	Relying on AI output and accepting output without assessment

Term	Your definition (In your own words)
Inside-frontier task	Complex tasks used in up to date and advanced LLMS
Outside-frontier task	Tasks that exceed what a LLM model is capable of doing

Self-Paced Tasks

Task 1: Professional Task Workflow Audit Matrix [30 min]

Purpose	Systematically assess five recurring tasks in your role for AI integration potential using evidence-based criteria, yielding a documented, reasoned selection of your top integration candidates.
----------------	---

[Click here to expand the text box](#)

Step-by-step instructions:

1. In the table below, list five tasks you perform at least monthly in your current role.
2. Rate each task on four dimensions (1–5 scale): (a) Time per instance, (b) Cognitive demand required, (c) AI integration potential (based on Module 2 knowledge), (d) Data sensitivity (1=Low, 5=High — reverse scored for integration).
3. For each task, rate Data Sensitivity: 1=No personal/confidential data; 3=Aggregated/de-identified data; 5=Student PII, FERPA, HIPAA, or proprietary research data.
4. Using the scores, identify your top two candidates (highest integration potential + lowest data sensitivity).
5. Write 2–3 sentences per selected candidate explaining your selection reasoning, referencing at least one Dell'Acqua et al. concept (jagged frontier, centaur/cyborg).

Workflow Audit Matrix:

DB:

In my Workflow Audit Matrix, one of my highest data-sensitivity tasks is reviewing data in student clinical logs which describe patient age, Dx and date of visit etc. There are more than 20 logs per day per assigned student on average. Data is de-identified and there are no names in the logs, however, a log with a patient who has a rare condition can potentially be identifiable. Although this a time consuming task especially depending on how many students you are supervising, this does not have high AI integration potential due to both student and clinical information.

There is tension (friction) with increasing workloads and the desire to give efficient and substantive student feedback. The risk is the temptation to ‘pencil whip’ a submission and giving a grade with little or no feedback. Students in the clinical setting require supervision and at times this supervision can vary. This is a challenging part in academia. The result is faculty working evenings, weekends and holiday to ensure the students get the appropriate and expected instruction from faculty.

	Task name	Frequency	Time/Instance	Cognitive demand (1-5)	AI integration potential (1-5)		Data sensitivity (1-5)	Top candidate? Y/N
1.		Student Clinical Logs	Daily or weekly	90 minutes per student	5	3	5	N
2.	Review SOAP Notes	3 x per month	60 min each note	5	5		4	N
3.	Review DNP Projects as Chair	5 x per month	4 hours each	5	5		4	N
4.	Emails from students re project questions and updates	Daily	30 minutes each over days	3	2		5	N
5.	Tracking advisee progress to degree and	Monthly	30 minutes per student spread over 3 days	2	5		2	Y

	Task name	Frequency	Time/Instance	Cognitive demand (1-5)	AI integration potential (1-5)	Data sensitivity (1-5)	Top candidate? Y/N
	whether full or part time						
6.	Tracking Doctoral project students progress as Chair	Monthly	30 minutes per student spread over 3 days	2	5	2	Y

Guidance Hints

* Be honest about time estimates. 'Drafting committee minutes' sounds like 30 minutes, but if you include chasing attendees for notes, formatting, and revision cycles, it may be 3 hours spread over 3 days. Accurate estimates are necessary for calculating actual time savings.

* Cognitive demand rating: tasks that require original professional judgment (curriculum design decisions, mentoring conversations, research interpretation) should rate 4–5. Tasks that are primarily information assembly, formatting, or routine communication should rate 1–2. High- demand tasks are usually poor AI integration candidates for the core reasoning step, though AI may help with peripheral steps.

* Data Sensitivity is a gateway criterion: any task rated 4 or 5 should be removed from AI integration consideration unless your IT team has specifically approved the tool for that data classification. This is not optional — it is an institutional risk management requirement.

* If all five tasks rate high on data sensitivity, note this as a finding and choose tasks with the lowest sensitivity for Task 2. The finding itself (high data sensitivity across your work) is worth discussing in the synchronous session.

Common Misconceptions to Avoid	
* MISCONCEPTION: 'My most time-consuming task is automatically the best AI integration candidate.' CORRECTION: Time cost alone is insufficient. Dell'Acqua et al. (2023) show that performance degrades on outside-frontier tasks even when they are time-consuming. Task type (inside vs. outside frontier) matters as much as task duration.	
* MISCONCEPTION: 'AI will automate my task completely.' CORRECTION: Noy & Zhang (2023) found a 37% average time reduction — significant, but not elimination. The 'jagged frontier' concept (Dell'Acqua et al., 2023) explains why: AI performs well on frontier-inside components of a task and poorly on frontier-outside components. Most real professional tasks span both zones.	
* MISCONCEPTION: 'If I'm hesitant to use AI for a task, I should skip it.' CORRECTION: Hesitation is analytically valuable. It usually signals either (a) genuine ethical concern (data sensitivity, academic integrity) worth investigating — or (b) uncertainty about AI capability for that task type. Both are worth surfacing in your selection reasoning.	
Best Practices	
* Complete the matrix for all five tasks before selecting your top two. The selection decision is better made with comparative data than with sequential intuition.	
* Identify a 'practice' candidate — a low-stakes version of a recurring task where failure would not cause professional harm. Your first AI-assisted workflow should be tested on the practice version before a consequential one.	
Task 2: AI-Assisted Task Completion with Full Documentation [90 min]	
Purpose	Generate documented, first-hand evidence of how AI changes your workflow for a real professional task — including where it adds value and where human judgment and correction remain essential.

[Click here to expand the text box](#)

Step-by-step instructions:

1. Select one of your top candidates from Task 1. Confirm data sensitivity is Low or Medium — if High, choose a different task.
2. Record your normal (unassisted) completion time for this task type from memory.

Get assignment for Advisees from Program Director

Get assignment for Clinical Supervising students from Program Director

Email preceptors – welcome letter

Get assignment for Student DNP projects from Program Director (identify role as chair, co-chair or third committee member)

Meet with students and identify background and career goals

Verify PT or FT status

Verify program of study

Track Case study meeting attendance

Track case study SOAP note

Track midterm in person visit ins multiple different clinical sites

Track midterm grading complete in Exxat and D2L

Site visits with midterm check across the U.S. (Some students have 2 or more sites during a semester).

Track meetings with Mentees (informal)

3. Complete the task using your chosen AI tool. Document EVERY prompt in sequence in the log below.
4. For each prompt, record the output (summarized) and the edits you made — specifically WHAT you changed and WHY.
5. After completion, record actual time spent and rate output quality vs. your typical unassisted quality (1–5).
6. Complete the Disclosure Decision box: Would you disclose AI use in this output? To whom? How? Reference your institution's policy by name.
7. Save both the raw AI output and your final edited version as separate documents.

Task Completion Log:

Task description (specific — not 'email' but the exact type, audience, and purpose):

Faculty is the audience working with doctoral Nursing students

Program director has full editing privileges

Faculty have editing privileges for assigned students

Program director gets students listing

Task description (specific — not 'email' but the exact type, audience, and purpose):

Program director gets city and state where student lives
Faculty Get assignment for Advisees from Program Director
Faculty Get assignment for Clinical Supervising students from Program Director
Faculty Email preceptors – welcome letter
Faculty Get assignment for Student DNP projects from Program Director (identify role as chair, co-chair or third committee member)
Faculty Meet with students and identify background and career goals
Faculty Verify PT or FT status
Faculty verify Cohort year
Faculty verify if degree seeking or certification
Faculty identify graduation year
Faculty identify first clinical semester and year
Faculty Verify program of study
Faculty Track Case study meeting attendance
Faculty determine if student is in a special grant or scholarship program
Faculty Track case study SOAP note
Faculty Track midterm in person visit ins multiple different clinical sites
T Faculty rack midterm grading complete in Exxat and D2L
Faculty Site visits with midterm check across the U.S. (Some students have 2 or more sites during a semester).
Faculty Track meetings with Mentees (informal)

Need reports and charts to track assigned student progress through degree program

Tool(s) used. Data sensitivity confirmed: Low / Medium:

Claude excel through U of A credentials

Normal time for this task (from memory): ____ . Actual time this session: ____

5 hours per week

PROMPT SEQUENCE — add rows for each prompt used:

Guidance Hints

* Document EVERY prompt, including failed or abandoned ones. The failed prompts are where the most learning happens — a prompt that produced an unusable output, followed by your diagnostic reasoning, is the most valuable entry in your log.

* The 'edits made' column is the most important. What you changed in the AI output reveals what the model got wrong or what it failed to capture about YOUR specific context, voice, or standards. If you made no edits, either the output was genuinely excellent or you accepted it too quickly — both conclusions deserve a sentence of reasoning.

* The 90-minute budget is intentional. This is not a speed exercise. Thorough documentation of a single task is more valuable than rushed documentation of two tasks.

* Reference your institution's AI use policy by name in the Disclosure Decision section. If your institution has no policy, write 'No current institutional policy exists' — this is itself a finding that AI Champions need to document and potentially address.

Common Misconceptions to Avoid

* MISCONCEPTION: 'If AI does most of the work, I should feel uncomfortable using it.'

CORRECTION: Professional tool use is not a measure of intellectual effort or moral worth. What matters is: (a) the output meets your professional standards, (b) it is disclosed appropriately, and (c) you take full responsibility for it. Noy & Zhang (2023) showed that AI-assisted professionals produced higher-quality outputs — quality, not effort, is the relevant standard.

* MISCONCEPTION: 'A task that took longer with AI than without AI means AI is not useful for that task.'

CORRECTION: Dell'Acqua et al. (2023) note that productivity gains emerge after participants develop prompting fluency. First-attempt AI-assisted workflows typically take longer than habitual unassisted workflows. Factor in learning curve costs when evaluating your first attempt.

* MISCONCEPTION: 'AI-generated content is the AI's responsibility, not mine.'

CORRECTION: All institutional outputs bear the full responsibility of the person who authorized and submitted them. AI is a tool. You are responsible for any output that carries your name or institutional affiliation.

Best Practices

* Save both the raw AI output and your final edited version. The before-and-after comparison is your primary evidence artifact for this module's self-assessment and is referenced in the Module 6 initiative design.

Best Practices

* Run the task for real — not a simulated version. Simulated tasks underestimate workflow complexity and produce different cognitive engagement than real professional stakes.

First looked at REDCap and SmartSheet to see what is available for faculty. Then went to Claude

Claude

	Prompts	Notes
1.	You are administration and faculty users that need a multifunctional database to identify student milestones. Create a database that has different faculty and different levels of access. Administrative users cannot assign faculty. Faculty is the audience working with doctoral Nursing students. Reports and charts are needed to track assigned student progress through degree program visually and through data. Create an excel database that includes: Program director has full editing privileges Faculty have editing privileges for assigned students Program director gets students listing Program director gets city and state where student lives Faculty Get assignment for Advisees from Program Director Faculty Get assignment for Clinical Supervising students from Program Director Faculty Email preceptors – welcome letter Faculty Get assignment for Student DNP projects from Program Director (identify rolle as chair, co-chair or third committee member) Faculty Meet with students and identify background and career goals Faculty Verify PT or FT status Faculty verify Cohort year Faculty verify if degree seeking or certification	This was a lot of prompts into one prompt and Claude actually did a really good job at ‘understanding’ what I needed. I was alarmed when Claude include test data. I had to comb through the excel to see the functionality.

	Faculty identify graduation year Faculty identify first clinical semester and year Faculty Verify program of study Faculty Track Case study meeting attendance Faculty determine if student is in a special grant or scholarship program Faculty Track case study SOAP note Faculty Track midterm in person visit ins multiple different clinical sites T Faculty rack midterm grading complete in Exxat and D2L Faculty Site visits with midterm check across the U.S. (Some students have 2 or more sites during a semester). Faculty Track meetings with Mentees (informal)	
2.	On each page with the student name, also include a column for student email. On the case study attendance page: Allow for 3 case study sessions for each semester Change column that states 'Attended' to 'Attended and participated' On the SOAP Note tracking page: Remove SOAP Note submitted column Change 'Faculty Reviewed' to 'Faculty graded in D2L' In midterm grading status page: Add column for 'Midterm graded in D2L & Exxat' Under Mentee meeting tab, allow spaces for multiple meetings per student per semester On Dashboard page allow the ability to filter by each faculty as well as cohort	Claude stated these were substantial changes.
3.	Please modify:	Locked out of Claude free version – Upgraded to continue – still saw limit.

	On Assignment page, add column for 'DNP Project Topic/Specialty' On student Program Info page, add a colum for 'Notes' On preceptor outreach page, allow the ability to review clinical sites and the ability to select Spring , Summer and fall semesters and what sites the student attended previously On the preceptor outreach page, allow space for up to 3 clinical sites per semester On the preceptor outreach page, allow site address column On midterm site visit page, allow site name, city, state and preceptor information auto populate on that page	Troubleshooted and logged out and logged back in with verifications etc.
4.		
5.		
6.		

For Discussion:

- What did AI do well:** I used Claude to create a student milestone tracker (similar to @Brett-Ratchford) but had a administrative and faculty focus. We have varying student assignments as clinical supervisor, advisor, project chair, etc. We also need to track grading milestones and interactions with students and how they are progressing through their doctoral projects. *I also looked at SmartSheet and REDcap as alternatives*). Claude seemingly intuitively created separate sheets for what was being tracked and include sample data so I was better able to see what was missing and what needed to be fine-tuned.
- What did it get wrong:** I had to prompt more to obtain more data fields for those students who have multiple clinical sites etc. It was my issue in prompting and not something Claude got wrong.
- What specific edit you made to correct or improve it.** I prompted Claude to make minor verbiage changes for students attending 'and participating' in online case study sessions, removed an assignment 'submitted' verbiage to assignment 'graded in D2L. I also added a new data field for Midterm graded in D2L & Exxat' and allowing free text fields, and allowing for tracking multiple clinical sites.

- **Overall:** this Task 2 allowed me to practice with Claude's capabilities and also establish a baseline tracker that our program may potentially use with necessary approvals.

Reply: Hi @randia-AZ,

I have prompted AI for ideas for questions as well. I have to specify that each answer choice must be plausible and the same length to avoid the issue you described. I also ask for 20 questions to get ideas and I can evaluate how I may be able to mix and match ideas to create a better questions. I have been an item writer for board certification exams and AI usage is strictly prohibited. The rule of thumb is that it takes 2 hours to create a good and unique exam question. The exam stats are crucial to determine if the questions assessed what it was intended to assess. It also still takes a great deal of time and expertise to verify. This was an important factor in this module's readings. Some of our tasks are not 100% conducive to a quick AI prompt especially depending on the subject matter, the criticality of accuracy (i.e. healthcare treatments) and how quickly available information (and guidelines) change which may exceed the model's capacity.

Task 3: Academic Integrity Scenario Analysis [30 min]

Purpose	Apply disclosure reasoning to four contested AI use scenarios using policy frameworks — developing the judgment needed to advise colleagues on appropriate AI use in professional university contexts.
----------------	--

Click here to expand the text box

Step-by-step instructions:

1. Open the [four scenarios from the Module 3 course platform resource](#).
2. For each scenario, write your verdict in the form below BEFORE discussing with peers.
3. Reference at least one named policy source per verdict (your institution's policy, or [IEEE/ACM ethical use guidelines](#)).
4. Note where your institution's policy is silent on the scenario — gaps are important findings for AI Champions.

Did you find any scenario where your institution's current policy was clearly silent?

What would you need a policy to say to make that scenario navigable?

I believe all of the scenarios presented have some nuances that may not be able to be fully covered by a policy per se. There are so many variables at play and to write a policy involving the evolving and emerging AI issues would be challenging to say the least. For

the Personalized Student Feedback scenario, I believe there are ethical concerns from offering AI only a few sentences out of an essay which may not provide correct feedback and subtleties of the entire paper. I found this resource with additional articles concerning AI in grading: <https://kpcrossacademy.ua.edu/the-digital-ia-harnessing-ai-for-fair-and-efficient-grading>

I also reviewed this U of A governance document that reiterates AI's limited capacity to understand context. I found this guideline to provide a good review of pertinent issues and concerns. I am happy to hear others' feedback and this discussion board comments. Link:: <https://emailarizona.sharepoint.com/sites/ISO-Communications/ISO%20Governance%20Documentation/Forms/AllItems.aspx?id=%2Fsites%2FISO%2DCommunications%2FISO%20Governance%20Documentation%2FISO%2D1500%20Vendor%20and%20Contract%20Security%2FISO%2D1500%2DG4%2DInformation%2DSecurity%2Dand%2DAI%2DTechnologies%2DGuideline%2Epdf&parent=%2Fsites%2FISO%2DCommunications%2FISO%20Governance%20Documentation%2FISO%2D1500%20Vendor%20and%20Contract%20Security>

Scenario Verdict Form:

Scenario	Appropriate? Y/N/Depends	Disclosure? Y/N	To whom	Policy cited	1-sentence reasoning
1 Personalized Student Feedback	N	Y	Student		Offering AI a few sentences of an essay may not provide correct feedback and subtleties. See https://kpcrossacademy.ua.edu/the-digital-ia-harnessing-ai-for-fair-and-efficient-grading AI limited ability to understand context – See https://emailarizona.sharepoint.com/sites/ISO-Communications/ISO%20Governance%20Documentation/Forms/AllItems.aspx?id=%2Fsites%2FISO%2DCommunications%2FISO%20Governance%20Documentation%2FISO%2D1500%20Vendor%20and%20Contract%20Security%2FISO%2D1500%2DG4%2DInformation%2DSecurity%2Dand%2DAI%2DTechnologies%2DGuideline%2Epdf&parent=%2Fsites%2FISO%2DCommunications%2FISO%20Governance%20Documentation%2FISO%2D1500%20Vendor%20and%20Contract%20Security

Scenario	Appropriate? Y/N/Depends	Disclosure? Y/N	To whom	Policy cited	1-sentence reasoning
					%20Governance%20Documentation%20ISO%201500%20Vendor%20and%20Contract%20Security
2 The Department Newsletter	Depends	Y	Recipients		
3 Promotion and Tenure Reference Letter	Depends	Maybe	Colleague and review committee	Ethics and policy	Most submissions for LOR specifically state cannot be AI generated. If used, must not include names or other potentially identifying data and would ensure that data training is turned off.
4 Student Disability Accommodation Letters	N	Y	IT Office and University	IT data governance policy	Potential FERPA violation and HIPAA violations. The pressure to use AI is coming from workload increases.

Guidance Hints

* Each scenario is deliberately ambiguous. There are no single correct verdicts. The goal is reasoned judgment under uncertainty, not rule application. The quality of your reasoning and the specificity of your policy citation matter more than the verdict itself.

* Distinguish between '**ethically questionable**' and '**policy-violating**.' These are not the same. Something can violate no current policy while still being ethically problematic (e.g., AI-drafted personalized feedback presented as wholly the author's own voice). Conversely, a policy can prohibit things that are ethically defensible.

Guidance Hints

* For scenarios involving student data: always ask whether the data was de-identified before AI processing, whether the tool was institutionally approved for data of that type, and whether FERPA obligations are met. These three questions apply regardless of policy statement clarity.

* When 'To whom' is unclear, think about who has a reasonable interest in knowing AI was involved: the recipient (if accuracy is important to them), the institution (if professional accountability is at stake), or the student (if the output directly affects them).

Common Misconceptions to Avoid

* MISCONCEPTION: 'Disclosure is only required when AI writes the whole thing.'

CORRECTION: Most emerging institutional and professional norms require disclosure of substantive AI assistance regardless of proportion. The relevant question is: **Would a reasonable person — knowing AI was used — evaluate this output differently? If yes, disclosure is appropriate.**

* MISCONCEPTION: 'If my institution has no AI policy, any use is implicitly permitted.'

CORRECTION: **Absence of policy is not permission.** FERPA, HIPAA, professional ethical norms, and general standards of honesty apply regardless of whether your institution has published an AI-specific policy.

* MISCONCEPTION: 'Disclosing AI use will damage my professional reputation.'

CORRECTION: The reputational risk is moving in the other direction: as AI detection practices improve and norms solidify, undisclosed AI use in professional contexts carries increasing reputational risk. **Proactive transparency is the norm for AI Champions.**

Best practices

* Write your verdict before discussing with peers. The peer exchange is most valuable when your initial positions differ and you must articulate the source of the disagreement — not when you pre-align to avoid conflict.

* **If your reasoning for two scenarios contradicts each other (you apply different standards to formally similar cases), examine that inconsistency. It usually reveals an unarticulated value difference worth surfacing.**

Task 4: Personal AI Workflow Sketch [30 min]

Purpose	Translate Task 2's documented AI-assisted task completion into a reusable, communicable workflow design — the foundation for the Module 6 initiative canvas.
----------------	--

[Click here to expand the text box](#)

Step-by-step instructions:

1. Using your Task 2 documentation as source material, complete the seven-box workflow canvas below.
2. Each box represents one decision point in your workflow. Be specific — name the tool, the data type, the check.
3. Identify at least ONE step where human verification is non-negotiable, regardless of AI output quality.
4. Add a brief note on the 'centaur vs. cyborg' style (Dell'Acqua et al.) that best describes your workflow.

AI Workflow Canvas — complete all seven boxes:

BOX 1 — Task entry point: What triggers this workflow? What is the desired output?
Program Director assignment
BOX 2 — Data inputs: What information goes in? Data sensitivity classification (Low/Med/High):
High sensitivity
Student names, grant status, city and state of residency
BOX 3 — AI tool selected + rationale: Why this tool for this task? Is it institutionally approved?
Claude used in interim, Requested REDCap access 07.03.2026.
BOX 4 — Prompt strategy: What prompting approach (zero-shot, few-shot, CoT)? Role specification?
CoT

BOX 5 — AI output review checkpoint: What specifically do you check? For accuracy? Tone? Bias? Completeness?

Accuracy in data needs and layout

BOX 6 — Human verification step (REQUIRED): What human action confirms quality before the output is used?

Checked every page with test data (provided by Claude which was very helpful). This allowed me to see what was missing and what needed to be reworded and / or omitted.

BOX 7 — Output + disclosure decision: Format, destination, and disclosure action (if any):

Excel format, would need to ensure we could use this tracker due to FERPA etc. Still need to look at REDCap, red booth, smart sheet as well with necessary U of A approvals.

Guidance Hints

* BOX 6 is where most first-draft workflows are weakest. 'Re-reading the output' is not a verification step — it is a fluency check. A verification step involves checking a specific claim against an external source, having a colleague review, or running a structured checklist.

* The 'centaur' style (Dell'Acqua et al.) means division of labor: you handle frontier-outside steps (complex judgment, relationship elements), AI handles frontier-inside steps (drafting structure, formatting, information assembly). The 'cyborg' style means deeper integration — AI embedded throughout. Note which applies to your workflow.

Common Misconceptions to Avoid

* MISCONCEPTION: 'A workflow that includes AI is automatically an AI workflow.'

CORRECTION: The workflow canvas is about intentional design — where AI sits, what it touches, and what remains human-controlled. Most good AI-integrated workflows are hybrids where AI handles specific steps and humans handle others. The design matters more than the proportion of AI involvement.

Best Practices

* Photograph or export your completed canvas. You will reference it in Module 6's Initiative Design task, where you will design an AI adoption workflow for colleagues.

In my Task 2 activity and in my professional role as well, I have the 'centaur' style (Dell'Acqua et al., 2023). As in our discussion reflected in our other discussion boards, I allow AI to assist with formatting, database creation, formatting and organization of information. I definitely oversee the relational aspects and use judgement and knowledge of human needs and human factors associated with any type of healthcare and healthcare student tasks, projects and responsibilities.

Peer Participation Guide & Self-Assessment Rubric

[Click here to expand the text box](#)

Peer Participation Guide

What to post in the course forum:

Post one concrete finding from your Task 2 Completion Log: what did AI do well, what did it get wrong, and what specific edit you made to correct or improve it. Do not include the full document — one specific example is enough.

How to respond to peers:

Read your peer's finding and respond with one sentence that connects their experience to either Noy & Zhang (2023) or Dell'Acqua et al. (2023). For example: 'Your experience with [specific finding] illustrates the jagged frontier concept — [specific connection].'

Discussion prompts to deepen the conversation:

- Looking at your Workflow Audit matrix: did your highest data-sensitivity tasks also have high AI integration potential? What does that tension mean for your institution?
- Did your Task 2 documentation show you operating as a 'centaur' or a 'cyborg' (Dell'Acqua et al.)? Which style do you think is more appropriate for your professional context — and why?
- In your Scenario Analysis, did you find any scenario where your institution's current policy was clearly silent? What would you need a policy to say to make that scenario navigable?

Forum etiquette for this module:

- Do not share full documents from Task 2 in the forum — they may contain institutional or sensitive content. Share specific, anonymized examples only.

- When connecting a peer's experience to a paper's finding, cite the paper specifically (author, year, concept name) rather than paraphrasing vaguely. This models the reference practice AI Champions use.

Self-Assessment Rubric

Score yourself honestly after completing all tasks. Bring unresolved low scores to the next synchronous session.

Criterion	Developing (1–2)	Meeting Expectations (3–4)	Exceeding Expectations (5)
Workflow Audit Quality	Fewer than 5 tasks listed; data sensitivity not assessed; selection reasoning absent or purely intuitive.	5 tasks listed with all dimensions rated; top 2 candidates selected with written reasoning; data sensitivity appropriately applied as a constraint.	Selection reasoning explicitly references Dell'Acqua et al. (centaur/cyborg or jagged frontier); one task is identified as a 'practice' candidate; inconsistencies in the matrix are acknowledged.
Task Completion Documentation	Task completion log is incomplete (fewer than 3 prompts documented, or edits not recorded); disclosure decision absent.	All prompts documented with output summaries and edit rationales; time comparison recorded; disclosure decision completed with policy reference.	Raw output and final version saved as separate files; disclosure decision references a named institutional policy; reflective quality comparison connects to Noy & Zhang findings.
Scenario Analysis Reasoning	Fewer than 4 scenarios addressed; verdicts with no policy citations; reasoning ignores data privacy dimensions.	All 4 scenarios addressed with verdicts, at least one policy citation per verdict, and 1-sentence reasoning; at least one policy gap identified.	All verdicts consistently reasoned; policy gaps explicitly noted; at least one verdict distinguishes between ethical appropriateness and policy compliance.
Workflow Canvas Completeness	Fewer than 5 boxes completed; Box 6 (human verification) is absent or says only 're-read'; data sensitivity not classified.	All 7 boxes completed with specific, not generic, content; Box 6 names a real verification action; centaur/cyborg style noted.	Canvas is specific enough that a colleague could follow the workflow independently; Box 6 names a concrete external verification source or method; canvas references Module 2 prompt strategy by name.

